Competitive Pricing Review

Tender IMD000ETT26000004 — IMDA AI Assurance Framework
Senior Executive Briefing | PwC Risk Services | 13 May 2026

★ **Insight**
What 'pricing' means in this tender: All non-volume bidders (37 of 42) submitted *average blended man-day rates*, not project totals. This means we are comparing **labour rates** — a much cleaner read on each firm's positioning than total-project price. The rate they bid is the rate they expect to charge — it reflects their seniority mix and target margin, not project scope.

Part 1 — Competitive Positioning

Executive Summary

Our footprint is intentionally narrow. PwC bid only 4 of 10 categories (2A, 2B, 3A, 3B) — the "do the testing / do the study" lots, not the "build the tools" lots. Coherent positioning, but it leaves **six categories where we are not in the conversation at all**.

Within our 4 categories, we are competitively priced — except 3B. We beat Deloitte and KPMG by 25–37% on 2A, 2B, and 3A. But on **3B (Technical Impact Studies) we are priced ABOVE Deloitte by 15%** ($368 vs $319). That is the one clear vulnerability.

The competitive threat is not just the Big 4. Five non-Big-4 firms are pricing inside the Big 4 band ($215–$300) and have AI-specific credibility (FAR AI, EquiStamp, DynamoFL, Revaisor, Ensign). They will be evaluated head-to-head with us on technical scoring.

The undercutters are the GLCs and Indian/Vietnamese MNCs. ST Engineering, NCS, FPT, Coforge, CMC-APAC are 35–80% cheaper than us. Price-weighted procurement loses to them. Quality-weighted procurement makes them noise.

Big 4 Footprint (excluding EY)

Firm	# Categories Bid	Pattern
Deloitte	7 of 10	Broadest scope; aggressive in Category 1 builds
PwC	4 of 10	Narrow — testing and impact studies only
KPMG	3 of 10	Most conservative; only 2A, 2B, 3B at one flat rate

Read: Deloitte is going wide. KPMG is going deep on one lane. PwC sits in between — picky but moderately so.

Head-to-Head on PwC's 4 Categories

Category	PwC	KPMG	Deloitte	PwC vs Deloitte	PwC vs KPMG
2A AI Testing	$263.75	$391.43	$417.88	–37% ■	–33% ■
2B Frontier Risks	$305.00	$391.43	$417.88	–27% ■	–22% ■
3A Societal Studies	$248.50	—	$379.40	–35% ■	n/a
3B Technical Studies	$368.33	$391.43	$319.20	+15% ■	–6% ■

3B is the only category where Deloitte undercuts us. If the evaluator weights price ≥40% on 3B, we lose to Deloitte head-to-head.

Wins, Exposures, Win-Probability

Wins (defensible):

• **2A** — Cheapest Big 4 by a wide margin. Combined with our brand and government relationships, the lot we're most likely to win on the Big-4 axis.

• **3A** — Deloitte is our only Big-4 competitor here, and we're 35% cheaper. Strong positioning.

Exposed:

• **3B** — Deloitte cheaper. Temus ($202.60), NCS ($259.20), TÜV SÜD ($225) are *all* cheaper than us. Five credible competitors below our price.

• **2B** — Not the cheapest on the panel. TÜV SÜD at $225 and NCS at $259 sit below us, both with valid AI safety / certification narratives.

• **Categories 1A–1E + 2C** — Not in the conversation. If the buyer awards a multi-lot framework, we won't be on it for build/tooling work — Deloitte will.

Win-Probability Scorecard

Lot	Our Position	Win Probability	Why
2A	Strong	HIGH	Cheapest Big 4, strong brand
2B	Mid	Mid	TÜV SÜD niche, NCS coverage threat
3A	Strong	HIGH	Only Deloitte to compete, far cheaper
3B	Weak	LOW–MID	Out-priced by 5 credible firms incl. Deloitte

Part 2 — Highlighted Non-Big-4 Competitors

Five named-watch competitors outside the Big 4. Their pricing reveals their strategy.

Resaro Singapore — AI assurance pure-play

Bids: 1A, 1B, 1C, 1D, 1E, 2A (six lots). **Rate band:** $276–$299.
Read: A specialist, not a generalist. Bid Category 1 (build tools) and 2A (test AI systems) — the technical AI assurance lots — and conspicuously skipped 2B, 2C, 3A, 3B. They know what they are. Pricing at the lower end of Big-4 band — exactly where an AI specialist sits to undercut us on technical credibility. **Most dangerous specialist competitor for Category 1 work — but not in our four lots.**

Temus — Temasek / Thoughtworks-backed digital firm

Bids: 1A, 1B, 1C, 1E, 2A, 2C, 3B (seven lots). **Rate band:** $202–$270.
Read: Aggressive on coverage and price. 70% of lots, 20–30% below Big 4 across the board. **Direct competitor on 2A and 3B**, and on 3B they are at $202.60 — far below our $368.33. Temasek + Thoughtworks DNA gives a "GLC-flavoured tech" story that is hard to counter on price alone. **Key competitor.**

NCS — Singtel-owned GLC, government incumbent

Bids: 9 of 10 (skipped 3A only). **Rate:** $259.20 flat across all nine lots.
Read: Flat-rate procurement-simplification play — "no haggling, one rate, full coverage." Combined with GLC status and incumbent government relationships, NCS is the **default safe-choice competitor**. Sits slightly below us on 2A and 3B, well below us on 2B. Strategic threat especially on government procurement bias.

ST Engineering Mission Software & Services — Defence / tech GLC

Bids: Nine lots (no 3A). **Rate band:** $134–$162.
Read: Heavy undercut — pricing 45–60% below us. Can sustain it because of scale, strategic foothold ambitions, and GLC margin tolerance. **The price alone is not the threat — the combination of low price + national-security narrative + GLC status is.** If the buyer cares about Singaporean sovereign capability, ST Engineering wins.

TÜV SÜD PSB — German certification body, niche play

Bids: 2B, 3B only (two lots). **Rate:** $225 flat.
Read: Deliberate narrow play on the certification / standards lots — 2B Frontier Risks and 3B Technical Studies. This is precisely where their TIC (testing / inspection / certification) DNA gives them a credibility advantage we don't have. $225 sits 25–40% below us on both lots. **Direct threat on 2B and 3B specifically.**

Threat Summary vs PwC

Firm	2A	2B	3A	3B
Deloitte	Low (price)	Low (price)	Low (price)	HIGH (price + scope)
KPMG	Low (price)	Low (price)	n/a	Mod (similar price)
Resaro	n/a	n/a	n/a	n/a
Temus	Moderate	n/a	n/a	HIGH (price)
NCS	Moderate	Moderate	n/a	Moderate
ST Engineering	Moderate	Moderate	n/a	HIGH (price)
TÜV SÜD	n/a	HIGH (price + niche)	n/a	HIGH (price + niche)

3B is the most contested lot for us. Five firms below our price, two with specialist credibility narratives (TÜV SÜD certification, Temus tech-DNA).

Part 3 — The Big-4-Adjacent Tier

Firms pricing in the $200–$400 band (PwC's range, excluding EY). These are the ones we will be evaluated *against* on technical scoring, not just price.

Firm	Origin	Rate Band	Lots	Why they matter
FAR AI	US AI safety	$286 flat	8	Strongest brand-credibility competitor outside Big 4. AI safety community recognises them. Priced just below us.
EquiStamp	US (PBC)	$215 flat	7	AI safety eval infrastructure. Works with METR, Epoch AI, EU Commission. Lowest credible specialist price in band.
Ensign Infosecurity	SG (Large local)	$220–$283	10	Cybersecurity heritage + Temasek / ST Eng backing. Full-lot coverage. Cyber-flavoured testing edge.
DynamoFL / Dynamo AI	US	$241–$262	6	Privacy-AI specialist. YC-backed. Strong in regulated industries.
Revaisor Ltd	UK	$270 flat	8	UK AI governance specialist. Targets financial services. Same price as Temus across 8 lots.
NxGen Asia	SG (ST Telemedia group)	$511.20 flat	8	Priced ABOVE us. Premium positioning that doesn't quite fit their brand size.
Optimum Solutions	SG (MNC scale)	$134–$162	6	Between low-cost IT-services tier and our tier. 4,500-employee IT body shop pushing upmarket on AI.

What this tier tells us about the market

- **The Big 4 price band is crowded.** Eight non-Big-4 firms (FAR AI, EquiStamp, Ensign, DynamoFL, Revaisor, NxGen, Temus, Resaro) are priced credibly within $215–$300. Brand premium alone no longer justifies our rate — we need a quality story.
- **The AI safety specialist niche is concentrated.** FAR AI, EquiStamp, Apollo Research, Revaisor, SaferAI all play here. They will dominate scoring on technical AI-evaluation criteria. Our defence: enterprise delivery capability, programme management, regulatory advisory depth.
- **Singapore-local Tier-1 (Ensign, Temus, NCS, ST Eng) are pricing aggressively to defend home turf.** Combined with policy preference for local capability (TechSG, AI Verify Foundation), they have a structural advantage we should price against.

Part 4 — Pricing Trends, Anomalies & Things to Note

The Pricing Landscape — Six Tiers

Tier	Rate / day	Who's there	Implied positioning
Offshore body shop	$34–$55	CI&T;, FPT	Cost-led IT services
Mid-offshore	$85–$130	SG AI Safety Hub, Coforge, CMC-APAC, AvePoint, DeeepLabs, Knovel, Verve, Optimum	Global delivery centre pricing
Singapore SME tier	$130–$200	ST Eng, Terra Systems, AZAAS, Modeak	Local commercial rate
Credible AI specialist	$200–$310	PwC, NCS, Temus, Resaro, FAR AI, EquiStamp, DynamoFL, Revaisor, OneSleeve, Ensign	**The real fight tier**
Premium Big 4	$310–$500	KPMG, Deloitte, Evi Fuelle	Brand + senior staffing
Outlier high	$500–$2,500	NxGen, WPH Digital, Reversec, EY	Premium, not all justified

Big trend: The credible AI specialist tier is dense. Ten firms — including us, both major GLCs, two Temasek-adjacent locals, and four AI specialist foreigners — are priced inside a $100 spread. **The premium for "being PwC" is no longer pricing differentiation. It is only useful as a quality argument.**

Trend #1 — Flat-Rate Bidding is the Dominant Strategy

16 of 42 firms bid the same rate across every lot they entered. This is a deliberate procurement bet — "we will staff every lot with the same seniority mix; easier for the evaluator; easier for us."

Who didn't flat-rate: PwC, Deloitte, EY varied rates by lot — detailed bid economics with lot-specific staffing. Coforge varied rates ($89–$105) — Indian IT majors are sophisticated bidders. Ensign split into two tiers ($220 and $282.50) — recognised some lots need senior cyber-AI experts and others don't.

Trend #2 — Big-4-Adjacent Convergence at ~$260–$280

Nine firms within a $66 spread: NCS $259.20, Temus $258.80–$270, PwC $263.75, Revaisor $270, Resaro $276.02, FAR AI $286, Ensign (low) $220, DynamoFL $241–$262. This is the implicit market consensus on what a mid-senior AI assurance consultant costs in Singapore. Anyone above needs a quality story. Anyone below is commoditising or has a cost-structure advantage.

Trend #3 — Singapore Local Tier-1 Pricing Aggressively

ST Engineering avg ~$144. Knovel ~$120. NCS $259.20. Temus ~$252. Ensign ~$246. **Systematically below the Big 4** except where the lot is a brand-name play. These firms have the strongest "win on policy preference + price" combination. If procurement applies any local-preference weighting, they sweep.

Notable Anomalies

Modeak's $1.00 bids on 2C and 3A

Modeak bid **$200.83 on every other lot, then $1.00** on Categories 2C and 3A. Most likely a panel-positioning play — they want to be on the framework regardless of margin, possibly because they expect to upsell, or those lots have low demand. **Worth flagging to evaluators — does the $1 bid violate framework integrity?**

AIDX Tech's $3.33 million total

Bid as **full project estimates**, not as a rate card. Man-day rates ($212–$350) are reasonable; assumption of 420–2,900 man-days per lot is not. Betting the buyer wants a delivery partner, not a panel rate. **May be evaluated as non-compliant** with the tender format.

Deloitte's 50% price swing within Category 1

$456.50 on 1A and 1B, then **$222.60** on 1D — a 51% drop for what looks like similar build work. Either 1D is genuinely cheaper to staff, they have existing leaderboard capability at marginal cost, or it's strategic underpricing as a foot-in-the-door. **Deloitte is treating 1D as a beachhead.**

WPH Digital at $1,508 flat across 9 lots

A Singapore SME pricing at near-EY levels — the only Singapore SME in EY territory. Banking on their first-in-world ISO 42001 for oil & gas certification as differentiator. Either confident or a misread that will price them out of every lot.

Reversec's lone $2,000 bid on 2A

Singapore offensive-cybersecurity firm bid **one lot only, at the highest rate in the field**. "We only want red-teaming at our specialist rate." **Most disciplined niche bid in the tender.**

EY's $2,516 bid on 3A specifically

EY's other lots are $1,240–$1,480. Then 3A jumps to $2,516 — 70% above their next-highest. Implies senior policy / economist time. Could also be an effective "we don't really want this lot" price.

PwC's reversed 3A/3B asymmetry

Most firms price 3A (societal) higher than 3B (technical) — Deloitte $379 vs $319, EY $2,516 vs $1,240. **PwC inverted this:** $248.50 on 3A, $368.33 on 3B. Either our technical-studies team is more senior, or our societal team more junior. **Worth checking internally** — may reflect a staffing assumption that doesn't match the market.

Apollo Research's single $85,580 lump sum on 3B

The UK's most famous frontier-AI evaluation lab bid only one lot, as a lump sum, and **not** on the testing / red-teaming categories (2A, 2B) where they have the strongest reputation. Unexpected. Likely they don't want framework-rate commitments — "engage us for specific projects only."

Participation Patterns

Most-contested vs least-contested lots:

Lot	# Bidders	Read
1E Agent testing suite	27	Hot category — everyone wants to be in agent testing
2A AI Testing	28	The main event — everyone bids
1A Benchmarks	26	Foundational AI eval work, broad demand
1D Leaderboards	18	Narrower technical capability, fewer entrants
3A Societal Studies	11	**Sparse** — most firms lack policy / social-science depth
2C Community Testing	10	**Sparsest** — concept ill-defined; only confident firms bid

Strategic implication: Lots with FEWER bidders (3A, 2C) are easier to win on technical scoring alone. PwC's 3A bid sits in an 11-firm field with Deloitte the only Big 4 competitor — **that's a high-probability lot.**

Things to Note

- **AI safety non-profit pricing floor** — SG AI Safety Hub $85, SaferAI $160. Mission-priced, not market-priced. Distorts the bottom of the credible range.
- **42 bidders is a strong signal.** Government AI assurance tenders rarely attract this many. AI assurance is one of the hottest 2026 service categories. Future tenders will only get more crowded — lock in framework position now.
- **No consortium bids.** All 42 standalone. Reflects framework rate-card structure — no incentive to consortium.
- **Specialist credibility vs price decoupling.** AI safety specialists (Resaro, FAR AI, EquiStamp, DynamoFL) are pricing AT or BELOW the Big 4. Historically you'd expect a specialist premium. The reverse is now true — biggest structural shift in the data.
- **Bid format diversity is an evaluation red flag.** Three formats: rate card (37), volume × rate (2), lump sum (3). Buyer methodology for normalising these will materially affect outcomes.

Summary — Bottom Line

- The market has converged. Mid-senior AI consultants in Singapore cost ~$260–$280/day. The Big-4 premium has eroded.
- Flat-rate bidding is dominant. Most firms aren't modelling lot-by-lot — they're submitting blended panel rates.
- Specialists are pricing AT the Big 4, not above. Competitive intensity in AI assurance is unusually high.
- GLCs and Singapore Tier-1 are pricing aggressively to defend home turf — likely with policy-preference tailwinds.
- Anomalies (Modeak $1, AIDX $3.3M, WPH $1,508, Deloitte 1D swing) signal procurement gaming, not honest pricing. Worth flagging.
- PwC's 3B pricing is internally inconsistent with our own 3A. Either intentional and defensible, or a recalibration miss.
- Biggest evaluation risk is bid-format normalisation. How procurement compares apples-to-oranges will decide several lots.